

1 The algorithmic method of proving circuit lower bounds

Recall that “ $\text{AC}^0[m]$ circuits” can use AND gates, OR gates, and MOD_m gates, all with unbounded fan-in. There are constants and literals at the bottom. When m is not a power of a prime, the class $\text{AC}^0[m]$ is poorly understood. As mentioned previously in this course, it is an open problem to show $\text{NP} \not\subseteq \text{AC}^0[6]$.

On the bright side, there is a line of work showing that there are “somewhat explicit” functions that cannot be computed by small $\text{AC}^0[6]$ circuits and similar models. In particular, Murray and Williams showed $\text{NQP} \not\subseteq \text{ACC}$ [MW18], where NQP is nondeterministic quasipolynomial time and $\text{ACC} = \bigcup_m \text{AC}^0[m]$. The full proof that $\text{NQP} \not\subseteq \text{ACC}$ is beyond the scope of this course, but we will present some elements of the proof. At a high level, the proof that $\text{NQP} \not\subseteq \text{ACC}$ has two steps. The first step is a nontrivial *satisfiability algorithm* for $\text{AC}^0[m]$ circuits:

Theorem 1 (Nontrivial satisfiability algorithm for ACC). *For all constants $m, d \in \mathbb{N}$, there exists a constant $\varepsilon > 0$ such that the following holds. Given the description of a depth- d $\text{AC}_d^0[m]$ circuit $C: \{0, 1\}^n \rightarrow \{0, 1\}$ of size at most 2^{n^ε} , it is possible to determine whether C is satisfiable in time 2^{n-n^ε} .¹*

The second step (ignoring some technicalities) is to show that for any circuit class \mathcal{C} , if there is a nontrivial satisfiability algorithm with the parameters described above, then $\text{NQP} \not\subseteq \mathcal{C}$. The second step is very interesting, but we will focus on the first step (Theorem 1) in these lecture notes.

2 Depth reduction for ACC circuits

Recall that $\mathbb{Z}[x_1, \dots, x_n]$ is the set of n -variate polynomials with integer coefficients.

Definition 1 (L_1 norm of a polynomial). If $h \in \mathbb{Z}[x_1, \dots, x_n]$, then we define $L_1(h)$ to be the sum of the absolute values of the coefficients.

Definition 2 (SYM^+). We define $\text{SYM}^+[k]$ to be the class of functions $C: \{0, 1\}^n \rightarrow \{0, 1\}$ of the form $C(x) = g(h(x))$, where $h \in \mathbb{Z}[x_1, \dots, x_n]$ satisfies $\deg(h) \leq k$ and $L_1(h) \leq 2^k$. Note that h is multilinear without loss of generality. The function $g: \mathbb{Z} \rightarrow \{0, 1\}$ can be arbitrary, but we emphasize that it is a function of just one integer variable.

You can double check that each function in $\text{SYM}^+[k]$ can be computed by a “SYM of AND of literals,” where the AND gates have fan-in at most k and the SYM gate has fan-in at most $2^{O(k)}$. Consequently, each function in $\text{SYM}^+[k]$ can be computed by a TC_3^0 circuit of size $2^{O(k)}$. The following theorem is a key step in the proof of Theorem 1, as well as being interesting in its own right.

Theorem 2 (Simulating $\text{AC}^0[m]$ circuits using SYM^+ circuits). *Let $m, d \in \mathbb{N}$ be constants. If $C: \{0, 1\}^n \rightarrow \{0, 1\}$ is an $\text{AC}_d^0[m]$ circuit of size $S \geq n$, then $C \in \text{SYM}^+[\text{polylog } S]$.*

When d and m are growing parameters, the best bound known is $C \in \text{SYM}^+[(\log S)^{O(d \cdot s)}]$, where s is the number of distinct prime factors of m [CP19]. In these lecture notes, we assume d and m are constant for simplicity.

¹We assume a *random access* model of computation throughout these lecture notes.

2.1 Simulating MOD_m gates using MOD_p gates

Lemma 1. *Let $p, e \in \mathbb{N}$ be constants, where p is prime and $e \geq 1$. Then $\text{MOD}_{p^e} \in \text{AC}^0[p]$.*

Proof. We prove it by induction on e . The base case $e = 1$ is trivial. For the inductive step, let $e \geq 2$, let $x \in \{0, 1\}^n$, and let N be the Hamming weight of x . We claim that²

$$\text{MOD}_{p^e}(x) = \text{MOD}_p(x) \vee \text{MOD}_{p^{e-1}} \left(\bigwedge_{i \in S_1} x_i, \dots, \bigwedge_{i \in S_{\binom{n}{p}}} x_i \right), \quad (1)$$

where $S_1, S_2, \dots, S_{\binom{n}{p}}$ is an enumeration of all size- p subsets of $[n]$. If N is not a multiple of p , this is trivial: $\text{MOD}_{p^e}(x) = \text{MOD}_p(x) = 1$. Now assume N is a multiple of p . In this case, observe that

$$\binom{N}{p} = \frac{N \cdot (N-1) \cdots (N-p+1)}{p \cdot (p-1) \cdots 1}.$$

In both the numerator and the denominator, only the first term is a multiple of p . Therefore, the exponent of p in the prime factorization of $\binom{N}{p}$ is one less than the exponent of p in the prime factorization of N . That is, $p^e \mid N$ if and only if $p^{e-1} \mid \binom{N}{p}$. Eq. (1) follows. By induction, Eq. (1) shows $\text{MOD}_{p^e} \in \text{AC}^0[p]$; note that $\text{poly}(\binom{n}{p}) = \text{poly}(n)$ since p is a constant. \square

More generally, let m be an arbitrary positive integer, with prime factorization $m = p_1^{e_1} \cdot p_2^{e_2} \cdots p_s^{e_s}$. Then

$$\text{MOD}_m(x) = \text{MOD}_{p_1^{e_1}}(x) \vee \cdots \vee \text{MOD}_{p_s^{e_s}}(x).$$

Thus, we can simulate an $\text{AC}^0[m]$ circuit using AND gates, OR gates, MOD_{p_1} gates, MOD_{p_2} gates, \dots , and MOD_{p_s} gates. The depth blows up by a constant factor and the size blows up polynomially, assuming m is a constant.

2.2 Eliminating one layer of MOD_p gates

Lemma 2 (Modulus-amplifying polynomials). *For every $k \in \mathbb{N}$, there exists a polynomial $M_k \in \mathbb{Z}[x]$ such that $\deg(M_k) = O(k)$, $L_1(M_k) = 2^{O(k)}$, and for every $x \in \mathbb{Z}$ and every $p \in \mathbb{N}$,*

$$\begin{aligned} x \equiv 0 \pmod{p} &\implies M_k(x) \equiv 0 \pmod{p^k} \\ x \equiv 1 \pmod{p} &\implies M_k(x) \equiv 1 \pmod{p^k}. \end{aligned} \quad (2)$$

Proof. Define

$$M_k(x) = \sum_{i=0}^{k-1} \binom{2k-1}{i} \cdot x^{2k-1-i} \cdot (1-x)^i.$$

The degree and L_1 bounds are straightforward. Observe that $M_k(x)$ is a multiple of x^k , which proves Eq. (2). Now suppose $x \equiv 1 \pmod{p}$. Then $1-x$ is a multiple of p , so $(1-x)^i \equiv 0 \pmod{p^k}$ whenever $i \geq k$. Consequently,

$$\begin{aligned} M_k(x) &\equiv \sum_{i=0}^{2k-1} \binom{2k-1}{i} \cdot x^{2k-1-i} \cdot (1-x)^i \pmod{p^k} \\ &= (x+1-x)^{2k-1} && \text{by the binomial theorem} \\ &= 1. \end{aligned} \quad \square$$

²Recall that we defined $\text{MOD}_m(x) = 1 \iff x_1 + \cdots + x_n \not\equiv 0 \pmod{m}$, which is opposite to the way many sources define it.

Lemma 3 (SYM^+ can simulate $\text{SYM}^+ \circ \text{MOD}_p$). Let $n, k \in \mathbb{N}$, let p be prime, and let $C: \{0, 1\}^n \rightarrow \{0, 1\}$ be a formula consisting of variables feeding into MOD_p gates feeding into a $\text{SYM}^+[k]$ gate. Then $C \in \text{SYM}^+[O(k^3 \cdot p \cdot \log n)]$.

Proof. By introducing dummy variables if necessary, we can write

$$C(x) = g \left(\left(\sum_{i=1}^L c_i \prod_{j=1}^k \text{MOD}_p(x_{ij1}, \dots, x_{ij\ell}) \right) \bmod p^{k+2} \right),$$

where $g: \mathbb{Z} \rightarrow \{0, 1\}$, each $c_i \in \mathbb{Z}$, we have $\sum_{i=1}^L |c_i| \leq 2^k$, and $\ell \leq n$. (Reducing mod p^{k+2} doesn't destroy any information, because the sum lies between -2^k and 2^k .) Therefore,

$$\begin{aligned} C(x) &= g \left(\left(\sum_{i=1}^L c_i \bigwedge_{j=1}^k \left(\sum_{t=1}^{\ell} x_{ijt} \not\equiv 0 \pmod{p} \right) \right) \bmod p^{k+2} \right) && \text{by definition of } \text{MOD}_p \\ &= g \left(\left(\sum_{i=1}^L c_i \cdot \mathbb{1} \left[\prod_{j=1}^k \sum_{t=1}^{\ell} x_{ijt} \not\equiv 0 \pmod{p} \right] \right) \bmod p^{k+2} \right) && \text{because a product of nonzero elements is nonzero in any field, including } \mathbb{F}_p \\ &= g \left(\left(\sum_{i=1}^L c_i \cdot M_{k+2} \left(\left(\prod_{j=1}^k \sum_{t=1}^{\ell} x_{ijt} \right)^{p-1} \right) \right) \bmod p^{k+2} \right) && \text{by Fermat's little theorem and modulus amplification.} \end{aligned}$$

The expression above has the format of SYM^+ : first we apply a multivariate polynomial, and then we apply a univariate function ("reduce mod p^{k+2} , then apply g "). The degree of the polynomial is at most $k \cdot (p-1) \cdot \deg(M_{k+2}) = O(p \cdot k^2)$. The L_1 norm of this polynomial is at most $2^k \cdot L_1(M_{k+2}) \cdot (\ell^{k \cdot (p-1)})^{\deg(M_{k+2})} = n^{O(p \cdot k^3)}$. \square

2.3 Simulating the entire circuit

Proof sketch of Theorem 2. There are several steps, but none is too difficult, given the tools that we have developed.

1. Replace each MOD_m gate with AND gates, OR gates, MOD_{p_1} gates, \dots , and MOD_{p_s} gates, as described in Section 2.1.
2. Replace each AND/OR gate with a probabilistic polynomial over the field \mathbb{F}_2 with error $0.1/S$ and degree $\ell = O(\log S)$. Note that a degree- ℓ polynomial over \mathbb{F}_2 is a $\text{MOD}_2 \circ \text{AND}_\ell$ circuit, where the MOD_2 gate has fan-in at most $S^{O(\log S)}$. Let \mathcal{D} be the resulting distribution over circuits.
3. Independently sample $t = O(n)$ circuits $C_1, \dots, C_t \sim \mathcal{D}$ and set $C(x) = \text{MAJ}_t(C_1(x), \dots, C_t(x))$. By Hoeffding's inequality and the union bound over all $x \in \{0, 1\}^n$, there is some fixing of C_1, \dots, C_t such that C computes f . Note that each C_i consists of MOD_2 gates, MOD_{p_1} gates, MOD_{p_2} gates, \dots , MOD_{p_s} gates, and AND_ℓ gates (with literals and constants at the bottom).
4. By introducing dummy gates if necessary, we can ensure that *all gates at the same level are of the same type*. In other words, we can compute f using a circuit of the following form:

$$\text{MAJ}_t \circ (\text{MOD}_2 \circ \text{MOD}_{p_1} \circ \dots \circ \text{MOD}_{p_s} \circ \text{AND}_\ell)^{O(1)}.$$

5. Note that $\text{MAJ}_t \in \text{SYM}^+[\log t]$. We eliminate the layers underneath the SYM^+ gate one by one to get a $\text{SYM}^+[k]$ circuit. To handle MOD_p layers, we use Lemma 3. To handle AND_ℓ layers, we use the trivial fact $\text{SYM}^+[k] \circ \text{AND}_\ell \subseteq \text{SYM}^+[k \cdot \ell]$. Since the number of layers is $O(1)$, we get $f \in \text{SYM}^+[\text{polylog}(S)]$. \square

3 The satisfiability algorithm

Lemma 4 (Fast multipoint evaluation of multilinear polynomials). *Let $h: \{0,1\}^n \rightarrow \mathbb{Z}$ be a multilinear polynomial with integer coefficients. Given h , represented as a list of 2^n coefficients, it is possible to compute $h(x)$ for all $x \in \{0,1\}^n$ in time $2^n \cdot \text{poly}(n) \cdot \text{polylog}(L_1(h))$.*

Proof. Let $s = L_1(f)$. If $n = 0$, then the problem is trivial. Otherwise, there are polynomials h_0, h_1 such that

$$h(x_1, \dots, x_n) = h_0(x_2, \dots, x_n) + x_1 \cdot h_1(x_2, \dots, x_n). \quad (3)$$

Considered as lists of coefficients, h_0 and h_1 are simply the first half and the second half of h , respectively. In particular, $L_1(h_0) \leq s$ and $L_1(h_1) \leq s$. We recursively compute $h_0(x)$ and $h_1(x)$ for all $x \in \{0,1\}^{n-1}$, and then we use Eq. (3) to compute $h(x)$ for all $x \in \{0,1\}^n$. The time complexity of this algorithm, $T(n, s)$, satisfies

$$T(n, s) \leq 2T(n-1, s) + 2^n \cdot \text{poly}(n, \log s),$$

because $|h_0(x)| \leq s$ and $|h_1(x)| \leq s$ for all $x \in \{0,1\}^{n-1}$. This implies $T(n, s) \leq 2^n \cdot \text{poly}(n, \log s)$. \square

Proof sketch of Theorem 1. First, let us design a satisfiability algorithm that runs in time $2^{\text{polylog}(S)} + 2^n \cdot \text{poly}(n)$, given a circuit of size S where $n \leq S \leq 2^n$.

1. By Theorem 2, it is possible to write C in the form $C(x) = g(h(x))$ where $h: \{0,1\}^n \rightarrow \mathbb{Z}$ is a multilinear polynomial satisfying $L_1(h) \leq 2^{\text{polylog } S}$. We only showed that such a representation *exists*, but it turns out that it can be *constructed* in quasipolynomial ($2^{\text{polylog } S}$) time, if we represent h as a list of monomials with nonzero coefficients and we represent g as a list of $2^{\text{polylog}(S)}$ output values. We omit the proof that g and h can be efficiently constructed.
2. Rewrite h as a list of 2^n coefficients (many of which may be zero). This can be done in time $2^{\text{polylog } S} + 2^n \cdot \text{poly}(n)$.
3. Compute $h(x)$ for all $x \in \{0,1\}^n$. By Lemma 4, this can be done in time $2^n \cdot \text{poly}(n)$.
4. Check whether there is some $x \in \{0,1\}^n$ such that $g(h(x)) = 1$. This can be done in time $2^n \cdot \text{poly}(n)$.

Now we are ready to design a satisfiability algorithm with the parameters described in the theorem statement. Given an $\text{AC}_d^0[m]$ circuit $C: \{0,1\}^n \rightarrow \{0,1\}$ of size 2^{n^ε} , we define $C': \{0,1\}^{n-n^\varepsilon} \rightarrow \{0,1\}$ by the rule

$$C'(x) = \bigvee_{y \in \{0,1\}^{n^\varepsilon}} C(xy).$$

Then C is satisfiable if and only if C' is satisfiable. The circuit C' is an $\text{AC}_{d+1}^0[m]$ circuit of size 2^{2n^ε} , so using the algorithm described above, we can decide whether it is satisfiable in time $2^{\text{poly}(n^\varepsilon)} + 2^{n-n^\varepsilon} \cdot \text{poly}(n)$. The theorem follows by picking a small enough ε . \square

References

- [CP19] Shiteng Chen and Periklis A. Papakonstantinou. “Depth reduction for composites”. In: *SIAM J. Comput.* 48.2 (2019), pp. 668–686. ISSN: 0097-5397. DOI: [10.1137/17M1129672](https://doi.org/10.1137/17M1129672).
- [MW18] Cody Murray and Ryan Williams. “Circuit lower bounds for nondeterministic quasi-polytime: an easy witness lemma for NP and NQP”. In: *Proceedings of the 50th Annual Symposium on Theory of Computing (STOC)*. 2018, 890–901. DOI: [10.1145/3188745.3188910](https://doi.org/10.1145/3188745.3188910).