

CMSC 28100

Introduction to Complexity Theory


Autumn 2025

Instructor: William Hoza



TMs can simulate all “reasonable” machines

- We could add various bells and whistles to the basic TM model
 - Left-right-stationary Turing machines
 - Multi-tape Turing machines
 - A Turing machine with a two-dimensional tape
- None of these changes has any effect on the power of the model



The Church-Turing Thesis

- Let $Y \subseteq \{0, 1\}^*$

Church-Turing Thesis:

There exists an “algorithm” / “procedure” for figuring out whether a given string is in Y **if and only if** there exists a Turing machine that decides Y .

← Intuitive notion

← Mathematically precise notion

Turing machines vs. your laptop

- **OBJECTION:**

- “Each individual Turing machine can only solve **one** problem.
- My laptop is a **single** device that can run **arbitrary** computations.
- Therefore, Turing machines don’t properly model my laptop.”



Email machine??



Zoom machine?? GitHub?? Pose computer



Photoshop machine??

Code as data

- The response to this objection is based on the “code as data” idea
- A Turing machine M can be encoded as a binary string $\langle M \rangle$
- Plan: We will show how to simulate a Turing machine M , given its encoding $\langle M \rangle$

Universal Turing machines

Theorem: There exists a Turing machine U such that for every Turing machine M and every input $w \in \{0, 1\}^*$:

- If M accepts w , then U accepts $\langle M, w \rangle$.
 - If M rejects w , then U rejects $\langle M, w \rangle$.
 - If M loops on w , then U loops on $\langle M, w \rangle$.
-
- One super-algorithm that contains all other algorithms inside it!

Example: Exercise 3

M

	Symbols				
	0	1	_	#	\$
a	(a, _, R)	(b, _, R)	(c, _, R)	(d, _, R)	
b	(y, 0, R)	(b, 0, R)	(c, 1, R)	(d, #, R)	
c	(y, 1, R)	(b, 1, R)	(c, _, R)	(d, #, R)	
d	(y, #, R)	(c, #, L)	(b, #, L)	(a, 0, L)	
e					
f					



Download

Upload



$\langle M \rangle$

```
... {"a": {"0": ["a", "_", "R"], "1": ["b",
"_", "R"], "_": ["c", "_", "R"], "#": ["d",
_", "R"], "$": null, "&": null, "%": null,
"@": null}, "b": {"0": ["y", "0", "R"], "1":
["b", "0", "R"], "_": ["c", "1", "R"], "#":
["d", ...
```

Autograder Results

1) Inputs that are not edge cases (0/3)

Test Failed: 'Accept' != 'Reject'

- Accept

+ Reject

: Your Turing machine behaves incorrectly

2) Edge case: Strings of zeroes (0/0.5)

Test Failed: 'Timeout' != 'Reject'

- Timeout

+ Reject

: Your Turing machine behaves incorrectly

3) Edge case: Strings beginning with 1 (0/0.5)

$\approx U$

Universal Turing machines

Theorem: There exists a single Turing machine U such that for every Turing machine M and every input $w \in \{0, 1\}^*$:

- If M accepts w , then U accepts $\langle M, w \rangle$.
 - If M rejects w , then U rejects $\langle M, w \rangle$.
 - If M loops on w , then U loops on $\langle M, w \rangle$.
-
- To properly **prove** it, we need to clarify how $\langle M \rangle$ is defined

Encoding a Turing machine as a string

- To encode a Turing machine $M = (Q, q_0, q_{\text{accept}}, q_{\text{reject}}, \Sigma, \sqcup, \delta)$:
 - WLOG, $|Q| = |\Sigma| = 2^k$ for some $k \in \mathbb{N}$
 - WLOG, $Q = \{0, 1\}^k$, $q_0 = 0^k$, $q_{\text{accept}} = 1^k$, and $q_{\text{reject}} = 01^{k-1}$
 - Encode $b \in \Sigma$ as $\langle b \rangle \in \{0, 1\}^k$, with $\langle 0 \rangle = 0^k$, $\langle 1 \rangle = 10^{k-1}$, and $\langle \sqcup \rangle = 1^k$
 - Encode $(q, b, D) \in Q \times \Sigma \times \{L, R\}$ as $\langle q, b, d \rangle = q\langle b \rangle\langle D \rangle \in \{0, 1\}^{2k+1}$
 - Then $\langle M \rangle = 1^k 0 \langle \delta \rangle$, where $\langle \delta \rangle$ is the list of $\langle \delta(q, b) \rangle$ for all $(q, b) \in Q \times \Sigma$

Universal Turing machines


Theorem: There exists a single Turing machine U such that for every Turing machine M and every input $w \in \{0, 1\}^*$:

- If M accepts w , then U accepts $\langle M, w \rangle := \langle M \rangle w$.
- If M rejects w , then U rejects $\langle M, w \rangle$.
- If M loops on w , then U loops on $\langle M, w \rangle$.

- **Proof sketch:** Next two slides


Initializing the simulation

- U is given $\langle M, w \rangle = 1^k 0 \langle \delta \rangle w$
- Initialize a tape containing $q_0 = 0^k$
- Initialize a tape containing $\langle \delta \rangle$
 - Note: $|\langle \delta \rangle| = 2^{2k} \cdot (2k + 1)$. Can compute using binary counter
- Initialize a tape containing $\langle w_1 \rangle \langle w_2 \rangle \dots \langle w_n \rangle$
 - Note: $\langle w_i \rangle = w_i 0^{k-1}$



Advancing the simulation

- Until the simulation reaches a halt state:
 1. Find $\langle \delta(q, b_i) \rangle = \langle q', b', D \rangle$ within $\langle \delta \rangle$
 - Idea: Treat $q\langle b_i \rangle$ as a number N in binary
 - Use a binary counter to go to position $N \cdot (2k + 1)$
 2. Replace q with q' and replace $\langle b_i \rangle$ with $\langle b' \rangle$
 3. Move this head k cells in direction D



Interpretation of universal Turing machines

- One piece of “hardware” that can run arbitrary “software”
- It’s a general-purpose, **programmable** computer
- This is why you don’t need a separate laptop for each task
- If you want to build a computer from scratch in some post-apocalyptic future, then **your job is to build a universal Turing machine**



The Church-Turing Thesis

- Let $Y \subseteq \{0, 1\}^*$

Church-Turing Thesis:

There exists an “algorithm” / “procedure” for figuring out whether a given string is in Y **if and only if** there exists a Turing machine that decides Y .

← Intuitive notion


← Mathematically precise notion

Humans vs. technology

- **OBJECTION:** “The Turing machine model is based on paper-and-pencil computation. Maybe we can solve undecidable problems using advanced **science and technology!**”



Hypercomputers



- A **hypercomputer** is a hypothetical device that can solve some computational problem that cannot be solved by Turing machines, such as SELF-REJECTORS
- Could it be possible to build a hypercomputer?
- We could try using quantum physics, antimatter, black holes, dark energy, superconductors, wormholes, closed timelike curves, ...

The Physical Church-Turing Thesis

- Let $Y \subseteq \{0, 1\}^*$

Physical Church-Turing Thesis:

It is **physically possible to build a device** that decides Y if and only if there exists a Turing machine that decides Y .

The Physical Church-Turing Thesis

- The standard Church-Turing thesis is a **philosophical** statement
- The Physical Church-Turing thesis is a **scientific law**
- Conceivably, it could be **disproven** by future discoveries... but that would be very surprising
- Analogy: Second Law of Thermodynamics
- Analogy: Cannot travel faster than the speed of light

Which problems
can be solved
through computation?

What are Turing machines
capable of?

Which languages are decidable?

Contrived vs. natural

- $\text{SELF-REJECTORS} = \{\langle M \rangle : M \text{ is a self-rejecting Turing machine}\}$
- We proved that SELF-REJECTORS is **undecidable**
- **OBJECTION:** “SELF-REJECTORS seems like a very **contrived** example.”
- **RESPONSE:** There are other undecidable languages that are **natural/well-motivated/interesting!**

The halting problem



- **Informal problem statement:** Given a Turing machine M and an input w , determine whether M **halts** on w .
- **The same problem, formulated as a language:**
$$\text{HALT} = \{\langle M, w \rangle : M \text{ is a Turing machine that halts on input } w\}$$
- It's the problem of **identifying bugs** in someone else's code! 🐞

Attempting to decide HALT



- Given $\langle M, w \rangle$:
 1. Simulate M on w
 2. If it halts, accept
 3. Otherwise, reject

