

CMSC 28100

Introduction to  
**Complexity Theory**

Spring 2025

Instructor: William Hoza



# The nature of this course

- In this course, we will study
  - The **mathematical and philosophical foundations** of computer science
  - The **ultimate limits** of computation
- This course will give you powerful **conceptual tools** for **reasoning about computation**
- There will be very little programming
- Homework and exams will be primarily proof-based

# Who this course is designed for

- CS students, math students, and anyone who is curious
- Prerequisites:
  - Experience with mathematical [proofs](#)
  - CMSC 27200 or CMSC 27230 or CMSC 37000, or MATH 15900 or MATH 15910 or MATH 16300 or MATH 16310 or MATH 19900 or MATH 25500

# Who this course is designed for


- It's okay if you don't consider yourself "theory-oriented"
- You belong here
- It's my job to give you resources so you can learn and succeed
- I also consider it my job to persuade you that complexity theory is important, interesting, enlightening, fun, cool, and worthy of your attention

# Class participation

- Please ask questions!
  - “What does that notation mean?”
  - “I forget what a \_\_\_\_\_ is. Can you remind me?”
  - “How do we know \_\_\_\_\_?”
  - “I’m lost. Can you explain that again?”

# Textbook

- Classic
- Popular
- High-quality
- Not free 😞



# Assessment

- 28 homework exercises
  - Exercises 1-4 are due **Tuesday, April 1**
- Midterm exam in class on 4/23
- Final exam at the end of the quarter

# My office hours

- Mondays (starting next week), 9am to 11am, JCL 205
- Stop by! This is a great time for discussions
  - If you are **confused/curious** about something, I'll try to help you figure it out
  - If you are **stuck** on a homework exercise, I'll try to think of a good hint
  - If you have a **complaint**, I'll listen and try to make things better



# Teaching assistants

- Zelin Lv
  - Office hours: Fridays, 2pm to 3pm, JCL 205
- Yakov Shalunov
  - Office hours: Thursdays, 5pm to 6pm, JCL 205

# Technology

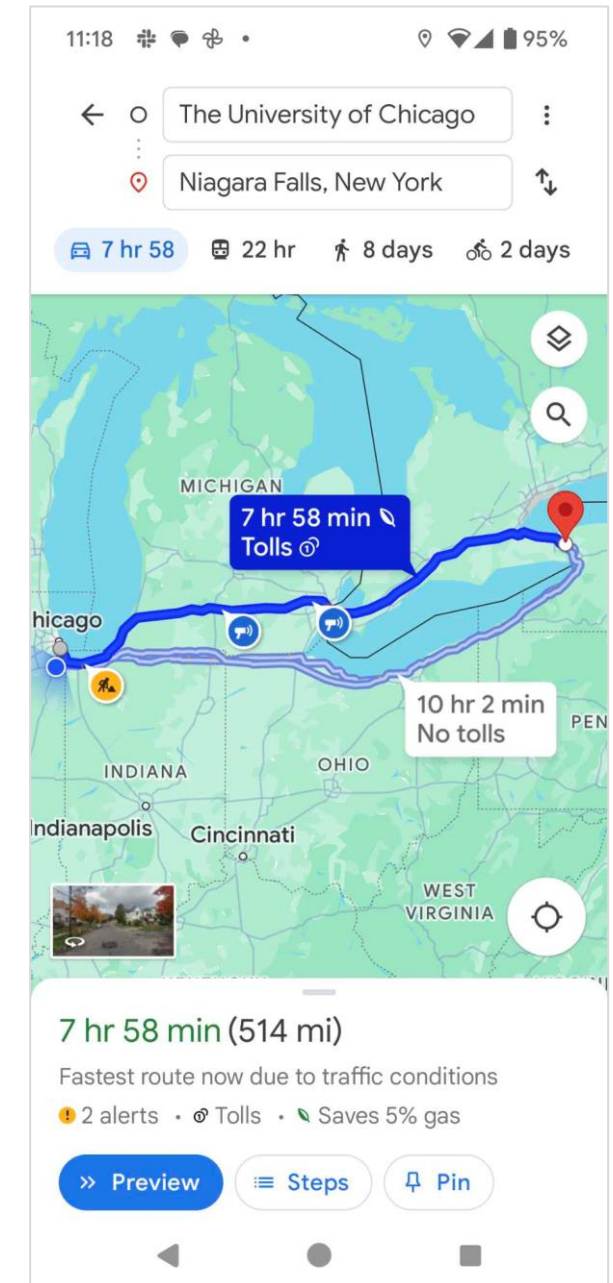
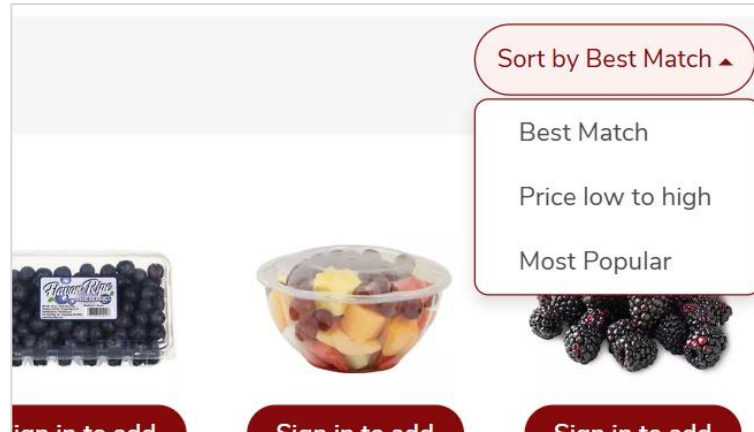
- Canvas: <https://canvas.uchicago.edu/courses/62607>
  - Homework exercises; practice exams; official solutions
- Course webpage: <https://williamhoza.com/teaching/spring2025-intro-to-complexity>
  - Course policies; slides
- Ed: <https://edstem.org/us/courses/76353/>
  - Discussions ( $\approx$  office hours); announcements
- Gradescope: <https://www.gradescope.com/courses/988815>
  - Submitting homework solutions; grades and feedback

The central question of this course:

**Which problems  
can be solved  
through computation?**

# Examples

- Many problems **can** be solved through computation:
  - Multiplication
  - Sorting
  - Shortest path
- Are there any problems that **cannot** be solved through computation?






# Impossibility proofs

- We will take a **mathematical** approach to this question
- We will formulate precise mathematical **models**
  - “Computation”
  - “Problem”
  - “Solve”
- Then we will write rigorous mathematical **proofs** of impossibility

Which problems  
can be solved  
through computation?

# Computation




- Computers: Modern technology?
- **Computation** is ancient
- Can be performed by:
  - A human being with paper and a pencil
  - A smartphone
  - A steam-powered machine
- We want a mathematical model that describes **all** of these and **transcends** any one technology





# Computation

- Note: Humans can do all the same computations that smartphones/laptops do
  - (less quickly and less reliably)
- Consequence: We can study computation without understanding electronics 😊
- Computation is a familiar, everyday, **human** act





# Ex: Palindromes

- Suppose a long string of bits is written on a blackboard

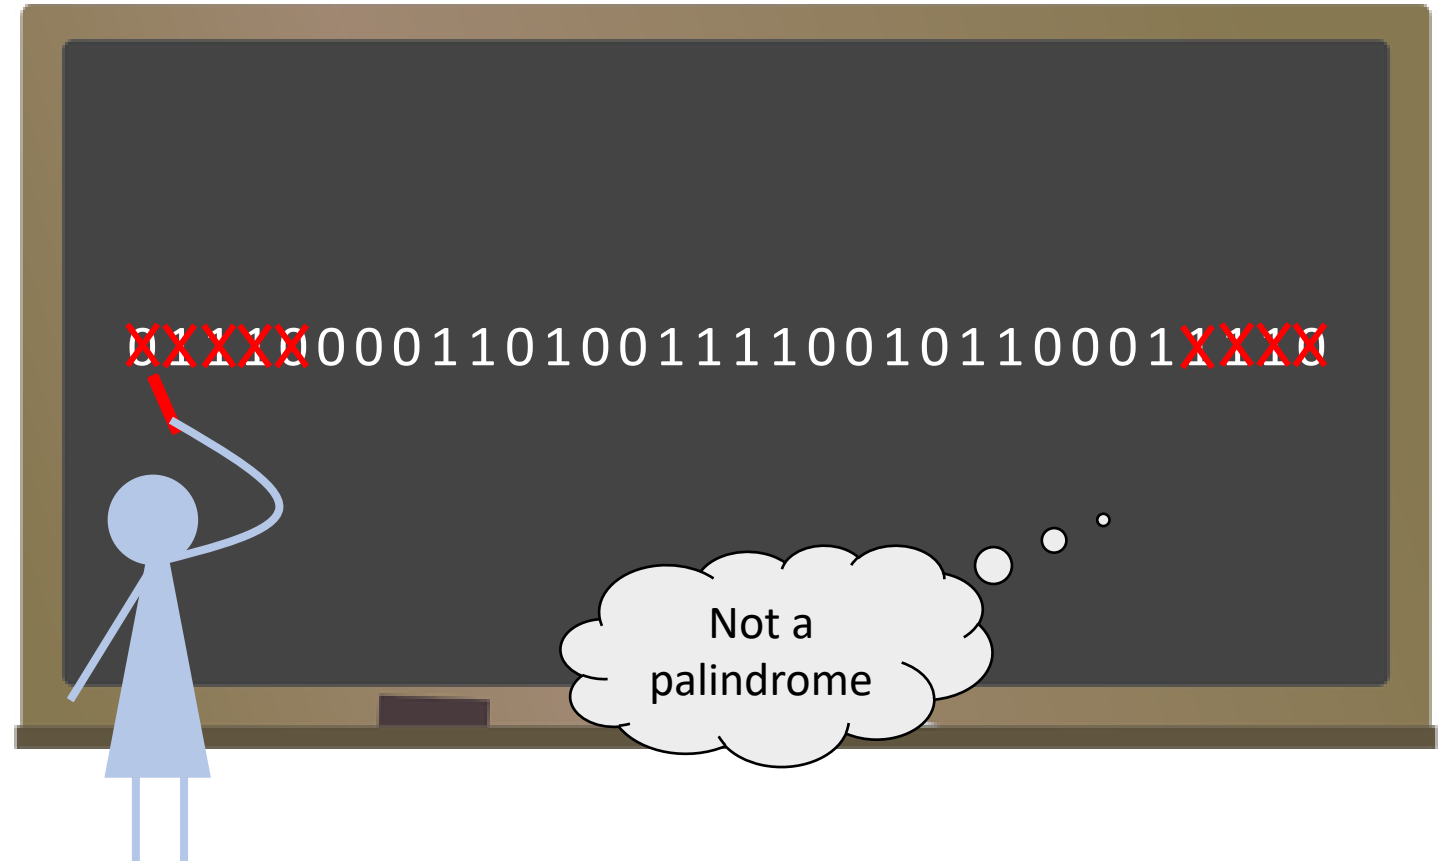
A blackboard with a dark grey surface and a brown frame. The string '01110000110100111100101100011110' is written in white chalk in the center. At the bottom of the frame, there is a small brown eraser on the left and a white piece of chalk on the right.

01110000110100111100101100011110

- Our job: Figure out whether the string is a “palindrome,” i.e., whether it is the same forwards and backwards
- What should we do?

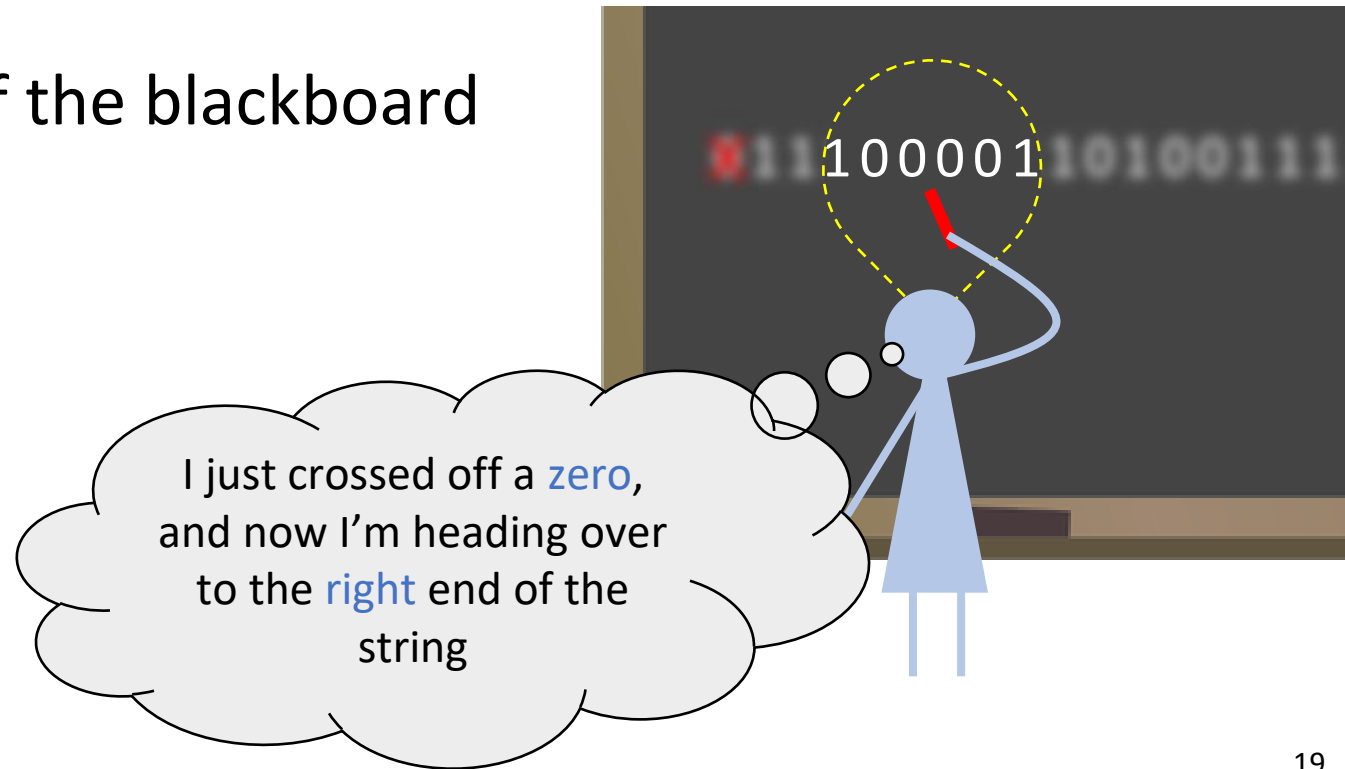
# Ex: Palindromes

- Idea: Compare and cross off the first and last symbols
- Repeat until we find a mismatch or everything is crossed off

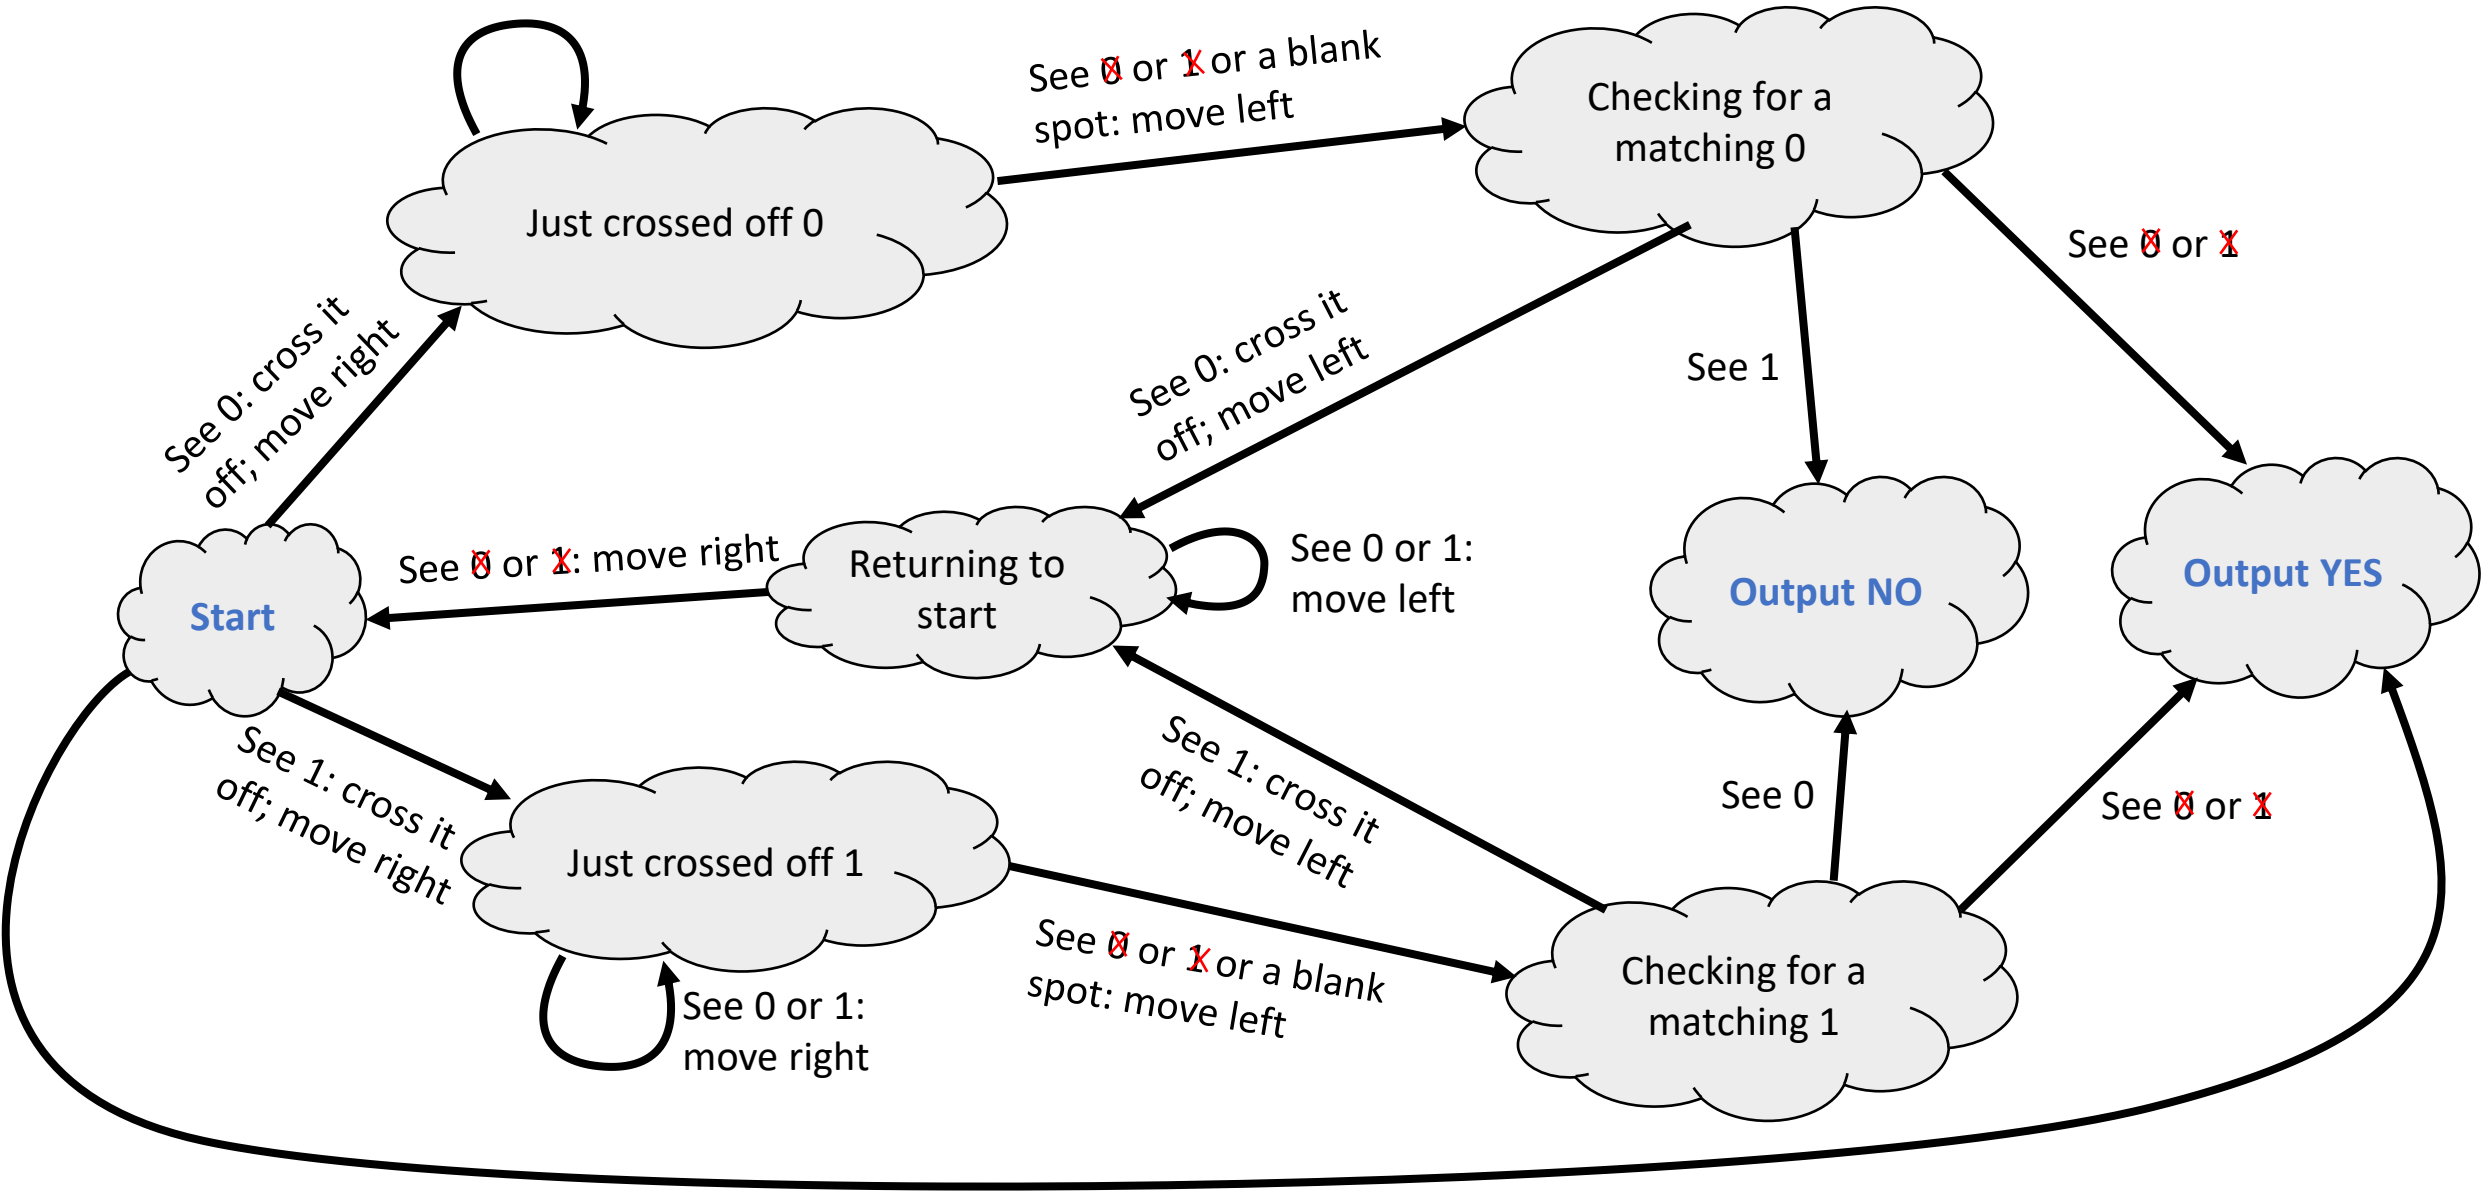


# Local decisions

- In each step, how do we know what to do next?
  1. We keep track of some information (“state”) in our **mind**
  2. We look at the **local** contents of the blackboard  
(one symbol is sufficient)
- We can describe the algorithm using a “**state diagram**”  
(next slide)



See 0 or 1: move right




See 0 or 1 or a blank spot

# The Turing machine model


- Turing machines: A **mathematical model of human computation**
- In a nutshell, a Turing machine is any algorithm that can be described by a **state diagram** like the one we just saw

# The Turing machine model



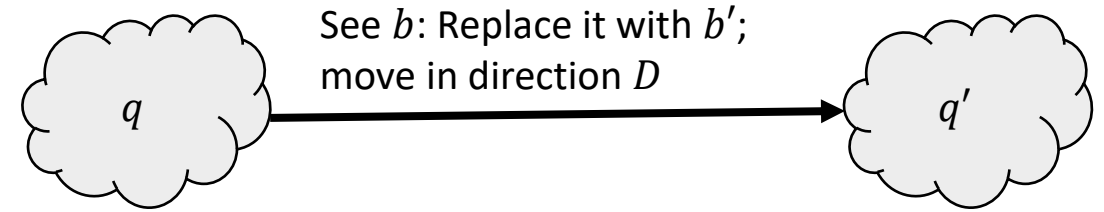
- We imagine an infinite, one-dimensional “tape”
- The tape is divided into “cells.” Each cell has one symbol written in it
- There is a “head” pointing at one cell of the tape
- The machine can be in one of finitely many internal “states”

# Turing machines



- In each step, the machine decides
  - What to write
  - Which direction to move the head (left or right)
  - The new state
- The decision is based only on the current state and the observed symbol

# Transition function



- Mathematically, we have a **transition function**

$$\delta: Q \times \Sigma \rightarrow Q \times \Sigma \times \{L, R\}$$

- Here  $Q$  is the set of **states** and  $\Sigma$  is the set of **symbols**
- $\delta(q, b) = (q', b', D)$  means:
  - If we are in the state  $q$  and we read the symbol  $b$ ...
  - Then our new state will be  $q'$ , we will write  $b'$  (replacing  $b$ ), and the head will move in direction  $D$ . (L = left, R = right)



# The input to a Turing machine

- One Turing machine represents one algorithm
- For us, the **input** to a Turing machine will always be a finite **string of bits**

# Symbols and alphabets

- An “alphabet”  $\Sigma$  is any nonempty, finite set of “symbols”
  - $\Sigma = \{0, 1\}$
  - $\Sigma = \{0, 1, \cancel{0}, \cancel{1}\}$
  - $\Sigma = \{A, B, C, \dots, Z\}$
  - $\Sigma = \{ \text{😍}, \text{⚾}, \text{🐍}, \text{🕒}, \text{🍕} \}$

# Strings

- Let  $\Sigma$  be an alphabet
- A **string** over  $\Sigma$  is a finite sequence of symbols from  $\Sigma$
- The **length** of a string  $x$  is the number of symbols, denoted  $|x|$
- If  $n$  is a nonnegative integer, then  $\Sigma^n$  is the set of length- $n$  strings over  $\Sigma$
- Example: If  $\Sigma = \{0, 1\}$ , then

$$\Sigma^3 = \{000, 001, 010, 011, 100, 101, 110, 111\}$$

If  $|\Sigma| = m$ , then what is  $|\Sigma^0|$ ?

A:  $|\Sigma^0| = 0$       B:  $|\Sigma^0| = m$

C:  $|\Sigma^0| = 1$       D:  $|\Sigma^0|$  is not well-defined

Respond at PollEv.com/whoza or text "whoza" to 22333

# The empty string

- If  $\Sigma$  is any alphabet, then  $|\Sigma^0| = 1$
- There is one string of length zero, called the **empty string**
- We use  $\epsilon$  to denote the empty string
  - Denoted `" "` in popular programming languages
- $\Sigma^0 = \{\epsilon\}$

# Arbitrary-length strings

- Let  $\Sigma$  be an alphabet
- We define  $\Sigma^*$  to be the set of strings over  $\Sigma$  of **any finite length**:


$$\Sigma^* = \bigcup_{n=0}^{\infty} \Sigma^n$$

- Example: If  $\Sigma = \{0, 1\}$ , then

$$\Sigma^* = \{\epsilon, 0, 1, 00, 01, 10, 11, 000, 001, 010, 011, \dots\}$$


# Turing machine initialization

- The tape initially contains the input string  $w \in \{0, 1\}^*$  (one bit per cell)
- Each cell to the left or right of the input initially contains a special “blank symbol”  $\sqcup$




# Turing machine initialization

- The head is initially at the cell containing the first bit of the input
- The machine is initially in a special “start state”  $q_0$




# Halting states



- There are two special “halting states,”  $q_{\text{accept}}$  and  $q_{\text{reject}}$
- If the machine ever reaches  $q_{\text{accept}}$ , this means it has **accepted** the input
- If the machine ever reaches  $q_{\text{reject}}$ , this means it has **rejected** the input
- Either way, the computation is finished. We say that the machine **halts**



# Looping




- It is also possible that the machine runs forever without ever reaching  $q_{\text{accept}}$  or  $q_{\text{reject}}$
- In this case, we say that the machine **does not halt**, does not accept the input, and does not reject the input

# Defining Turing machines rigorously

- **Definition:** A **Turing machine** is a 7-tuple  $M = (Q, q_0, q_{\text{accept}}, q_{\text{reject}}, \Sigma, \sqcup, \delta)$

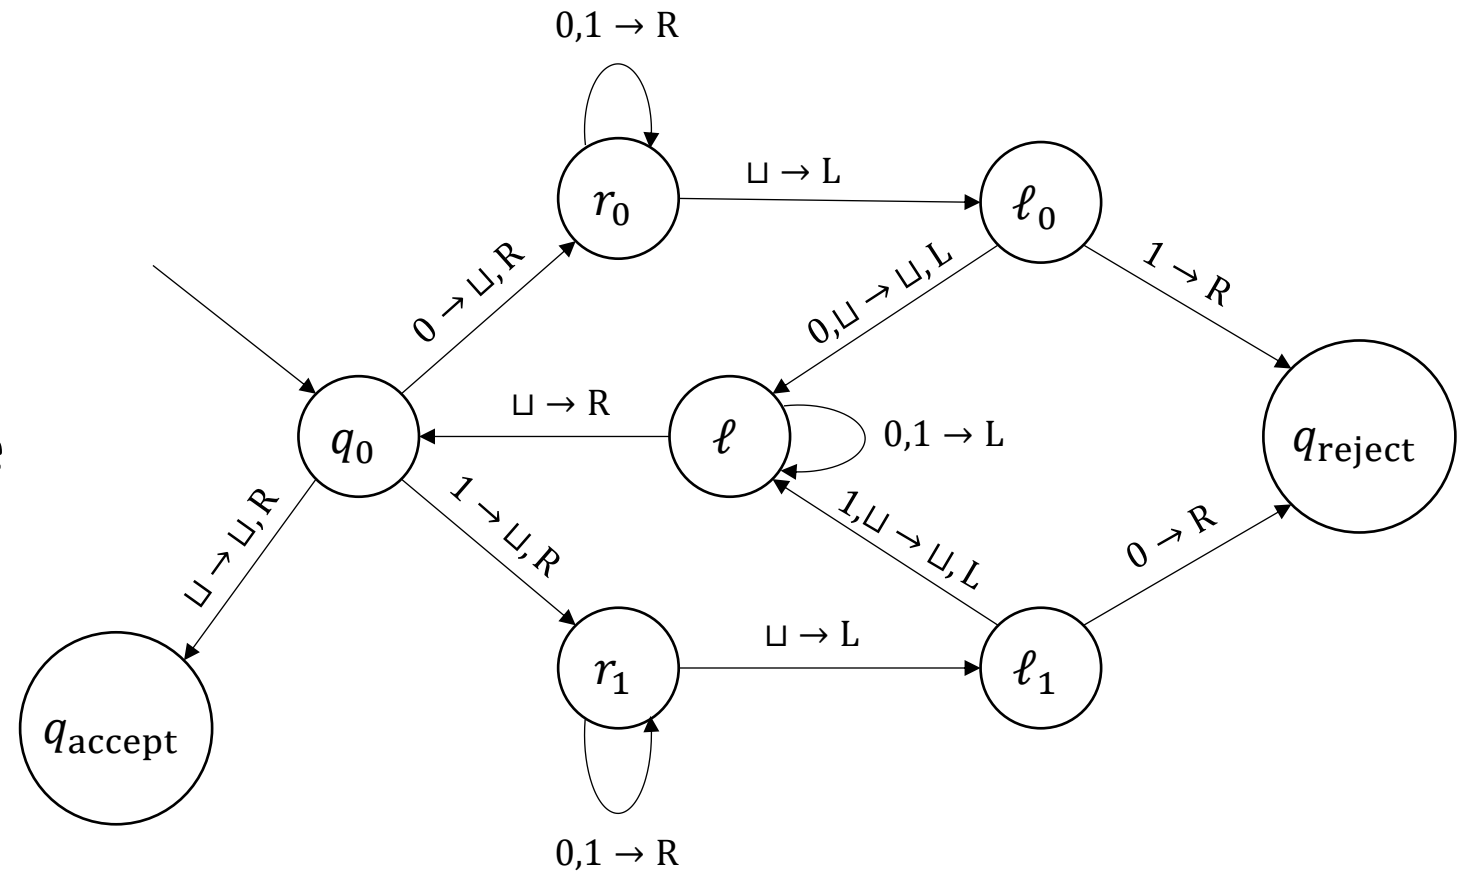
such that

- $Q$  is a finite set (the set of “states”)
- $q_0, q_{\text{accept}}, q_{\text{reject}} \in Q$  and  $q_{\text{accept}} \neq q_{\text{reject}}$
- $\Sigma$  is a finite set of symbols (the “tape alphabet”)
- $\sqcup$  is a symbol (the “blank symbol”)
- $\{0, 1, \sqcup\} \subseteq \Sigma$  and  $\sqcup \notin \{0, 1\}$
- $\delta$  is a function  $\delta: Q \times \Sigma \rightarrow Q \times \Sigma \times \{L, R\}$  (the “transition function”)

 Warning: The definition in the textbook is slightly different. Sorry!  
(The two models are equivalent.)

# State diagram

- Each node represents a state
- An arc from  $q$  to  $q'$  labeled “ $b \rightarrow b', D$ ” means  $\delta(q, b) = (q', b', D)$
- The label “ $b \rightarrow D$ ” is shorthand for “ $b \rightarrow b, D$ ”
- An arc labeled “ $a, b \rightarrow \dots$ ” represents two arcs (“ $a \rightarrow \dots$ ” and “ $b \rightarrow \dots$ ”)




# Defining TM computation rigorously

- Transition function  $\delta$  describes the **local** evolution of the computation
- What about the **global** evolution?

# Configurations of a Turing machine

- Let  $M = (Q, q_0, q_{\text{accept}}, q_{\text{reject}}, \Sigma, \sqcup, \delta)$  be a Turing machine
- A **configuration** of  $M$  is a triple  $(u, q, v)$  where  $u \in \Sigma^*$ ,  $q \in Q$ ,  $v \in \Sigma^*$ , and  $v \neq \epsilon$ . Interpretation:
  - The tape currently contains  $\cdots \sqcup \sqcup \sqcup \sqcup uv \sqcup \sqcup \sqcup \sqcup \cdots$
  - The machine is currently in state  $q$  and the head is pointing at the first symbol of  $v$



# Configuration shorthand

- Instead of  $(u, q, v)$ , we often write  $uqv$
- We think of  $uqv$  as a string over the alphabet  $\Sigma \cup Q$
- This shorthand can only be used if  $Q \cap \Sigma = \emptyset$ , which we can assume without loss of generality by renaming states if necessary

# Equivalent configurations

- Note:  $uqv$  and  $uqv \sqcup$  are technically two distinct configurations...
- However, they represent the **same scenario**
- We can say that they are **“equivalent”**
- (A configuration is a finite string, even though the tape is infinitely long)
- Similarly,  $\sqcup uqv$  is equivalent to  $uqv$

# The initial configuration

- Let  $w \in \{0, 1\}^*$  be an input
- The initial configuration of  $M$  on  $w$  is

$$\begin{cases} q_0 w & \text{if } w \neq \epsilon \\ q_0 \sqcup & \text{if } w = \epsilon \end{cases}$$



# The “next” configuration

- For any configuration  $uqv$ , we define  $\text{NEXT}(uqv)$  as follows:
  - Break  $uqv$  into individual symbols:  $uqv = u_1u_2 \dots u_{n-1}u_nqv_1v_2v_3 \dots v_m$
  - If  $\delta(q, v_1) = (q', b, \text{R})$ , then  $\text{NEXT}(uqv) = u_1u_2 \dots u_{n-1}u_nbq'v_2v_3 \dots v_m$ 
    - Edge case: If  $m = 1$ , then  $\text{NEXT}(uqv) = u_1u_2 \dots u_{n-1}u_nbq' \sqcup$
  - If  $\delta(q, v_1) = (q', b, \text{L})$ , then  $\text{NEXT}(uqv) = u_1u_2 \dots u_{n-1}q'u_nbv_2v_3 \dots v_m$ 
    - Edge case: If  $n = 0$ , then  $\text{NEXT}(uqv) = q' \sqcup b'v_2v_3 \dots v_m$
- We write  $\text{NEXT}_M(uqv)$  if  $M$  is not clear from context

# Halting configurations

- An **accepting configuration** is a configuration of the form  $uq_{\text{accept}}v$
- A **rejecting configuration** is a configuration of the form  $uq_{\text{reject}}v$
- A **halting configuration** is an accepting or rejecting configuration